

Renegade Whitepaper

Protocol Specification, v0.6

Christopher Bender
chris@renegade.fi

Joseph Kraut
joey@renegade.fi

Abstract

RENEGADE¹ is an on-chain dark pool. In contrast to other non-custodial decentralized exchanges, Renegade maintains complete anonymity during the entire lifecycle of a trade.

Peer-to-peer order matching is inferred via a secure multi-party computation, and atomic settlement of matched orders is performed via zero-knowledge proofs of valid matching engine execution. As a result, no third party (including the block proposer) can learn any information about any user's token balances, pending orders, or trade history, ultimately leading to minimal miner extractable value and higher-quality trade execution at midpoint prices.

1 Introduction

Currently, non-custodial trading suffers from four significant problems:

- **Miner Extractable Value.** Block producers (in a L1 context) and/or sequencers (in a L2 context) can see full transaction information, allowing for arbitrary reordering, frontrunning, backrunning, and trade censorship.
- **Pre-trade transparency.** Non-marketable trades that rest on a limit order book are visible to all third-parties, leading to quote fading.
- **Post-trade transparency.** All third-parties can query the entire state history, allowing for tracking and tracing of trading activities.
- **Address discrimination.** Traders can see the origination address (pseudonymous identity) of all outstanding orders, leading to worse fills against sophisticated counterparties.

All of these design flaws lead to worse trade execution, particularly for whale traders with large market impact.

BACKGROUND ON DARK POOLS.

A *dark pool* is a well-understood feature of traditional finance market structure. Legally classified as “alternative trading systems”, dark pools are off-exchange trading venues with better privacy protections for traders.

Dark pools have the same functionality as the more familiar “lit” exchanges like the NYSE or NASDAQ, with one

important difference: The order book is not publicly visible, meaning that traders cannot see the outstanding quotes of others. Participants in a dark pool are only informed about matches on their own trades, allowing for traders to privately search for a counterparty without broadcasting their trading intentions to the wider market.

Dark pools are typically used for better execution of trades on large blocks of equities. If a large trade were to be executed at once on a lit exchange, insufficient liquidity would lead to significant price impact and cross-exchange arbitrage. TWAP-style orders are often employed to help massage the order into the lit market over time, but statistical arbitrageurs can often detect such patterns, once again leading to quote fading and inferior execution.

In a crypto context, the block trade problem is even worse: Not only do current decentralized exchanges leak the current state of the order book, but blockchains inherently have fully-auditable state history. In addition to seeing the current state of the order book, any third-party can trace the past activity of all participants.

OUR SOLUTION.

In this paper, we introduce Renegade, a non-custodial dark pool. We solve all four problems outlined previously by hiding all information about the state of the exchange with zero-knowledge proofs.

Renegade is functionally equivalent to a CLOB-style DEX, but with an encrypted and distributed order book. Matches between users' orders are inferred via a cryptographically secure multi-party computation. Once a match has been found, settlements of swapped tokens are done via zero-knowledge proofs to hide all trade information while maintaining consistency of the system.

In order to implement this hybrid MPC-ZKP match-settle architecture, we use the **collaborative SNARK** framework from Ozdemir et al.² Essentially wrapping zero-knowledge proof generation inside of a MPC, collaborative proofs allow for traders to transact while leaking zero information to third-parties.

¹Renegade is hiring! Check out our [jobs page](#) and get in contact [on Twitter](#).

²Ozdemir and Boneh, *Experimenting with Collaborative zk-SNARKs: Zero-Knowledge Proofs for Distributed Secrets*, <https://eprint.iacr.org/2021/1530>

In the remainder of this paper, we give a formal protocol specification for Renegade.

We start by giving a general overview of the protocol in Section 2, defining the idea of a *wallet* that maintains private state. We describe how wallets are created and destroyed, showing how the *commitment tree* can allow full user state privacy. From this, we give a precise specification of all state that is maintained by the system, both on the client side and on the smart contract side.

Next, in Section 3, we illustrate the full lifecycle of a trade from wallet creation to trade settlement, and explain the core collaborative proving protocol that allows for completely anonymous trading.

Then, in Section 4, we describe our peer-to-peer protocol for counterparty discovery. We illustrate the network topology of the system, introducing the concept of a *relayer*. In addition, we define the key hierarchy that allows for various levels of access controls into a trader's wallet.

Next, in Section 5, we describe the concept of *indications of interest*, showing how Renegade allows for individual traders to trade off execution quality for execution latency.

Finally, in the Appendix A, we give formal specifications for the seven different NP statements that are used to ensure state consistency within Renegade.

2 Protocol Overview

In this section, we describe the stateful elements of the protocol. Later, in Section 3, we define the rules by which this state may be updated.

2.1 Wallets

Balance and order information cannot be recorded on-chain in-the-clear (or else privacy would be compromised), so traders instead maintain a private **wallet** containing all local state. A wallet has two primary functions:

- Keep track of what tokens a trader owns: The smart contract only maintains global pools of all token deposits, so we need to privately keep track of which trader owns what tokens.
- Keep track of the trader's set of outstanding orders: We cannot post order information transparently, so we keep the set of open orders inside of the wallet.

Note that the wallet W may be kept off-chain, or the wallet may be encrypted with a view keypair and written on-chain.³ The view keypair is deterministically derived from the user's native Ethereum keypair, so a trader only needs to have access to their Ethereum key in order to recover their wallet.

³For strongest guarantees against data loss, the user may choose to store the encrypted wallet as on-chain L1 calldata, paying high gas cost. However, if the user is particularly fee-sensitive, they may want to instead store the wallet in some off-chain data availability layer, avoiding the more expensive calldata fees.

Formally, let $M_O, M_B, M_F \in \mathbb{N}$ be public constants defining the maximum number of orders, balances, and fee approvals, respectively, that a user may have at once. Now, for some elliptic curve \mathbb{G} over a prime field \mathbb{F} , a wallet W is defined as a tuple

$$W := (B, O, F, K, r) \in \mathbb{F}^{2M_B+8M_O+5M_F+9}$$

with the following definitions:

- $B = (m_i, v_i)_{i \in [M_B]}$ is a list of size M_B of elements of \mathbb{F}^2 :
 - $m_i \in \mathbb{F}$ is the mint (i.e. contract) address of a token that is held in this wallet.
 - $b_i \in \mathbb{F}$ is the amount of this token that is held in this wallet.
- $O = (\lambda_i, m_{1_i}, m_{2_i}, s_i, p_i, a_i, \alpha_i, \tau_i)_{i \in [M_O]}$ is a list of size M_O of elements of \mathbb{F}^8 :
 - $\lambda_i \in \mathbb{F}$ is a flag that denotes the type of the order (0 is limit, 1 is midpoint-pegged, etc.).
 - $m_{1_i} \in \mathbb{F}$ is the mint address of the base token (e.g. WETH).
 - $m_{2_i} \in \mathbb{F}$ is the mint address of the quote token (e.g. USDC).
 - $s_i \in \mathbb{F}$ is the side of the order (0 is buy, 1 is sell).
 - $p_i \in \mathbb{F}$ is the limit price (in units of quote per base), encoded in fixed-point representation. Ignored if the order is not a limit type.
 - $a_i \in \mathbb{F}$ is the amount of base currency that the user wants to buy or sell.
 - $\alpha_i \in \mathbb{F}$ is the minimum fill size, in units of the base currency.
 - $\tau_i \in \mathbb{F}$ is the timestamp of when this order was last updated, and is used for limit order price improvement.
- $F = \left(\text{pk}_{\text{cluster}_i}^{\text{settle}}, \mu_i, \delta_i, \gamma_i \right)_{i \in [M_F]}$ is a list of size M_F of elements of $\mathbb{G} \times \mathbb{F}^3$:
 - $\text{pk}_{\text{cluster}_i}^{\text{settle}} \in \mathbb{G}$ is the public settle key of some relay cluster that is allowed to take fees for managing this wallet.
 - $\mu_i \in \mathbb{F}$ is the mint address of some token that is used to pay for gas fees.
 - $\delta_i \in \mathbb{F}$ is the constant amount of μ_i that this cluster may debit from the balances B upon any call to the contract.
 - $\gamma_i \in \mathbb{F}$ is the fixed-point percentage fee that the cluster may take upon any successful match.
- $K = \left(\text{pk}^{\text{root}}, \text{pk}^{\text{match}}, \text{pk}^{\text{settle}}, \text{pk}^{\text{view}} \right)$ is a 4-tuple of elements of \mathbb{G} :

- ▶ $pk^{\text{root}} \in \mathbb{G}$ is the public key that corresponds to a secret key $sk^{\text{root}} \in \mathbb{F}$ that must be known in order to deposit/withdraw from the balances B , or to update the order book O . This is typically a secret key controlled by the end user / trader.
- ▶ $pk^{\text{match}} \in \mathbb{G}$ is the public key that corresponds to a secret key $sk^{\text{match}} \in \mathbb{F}$ that must be known in order to match outstanding orders in the order book O . This is typically a secret key controlled by a relayer.
- ▶ $pk^{\text{settle}} \in \mathbb{G}$ is the public key that corresponds to a secret key $sk^{\text{settle}} \in \mathbb{F}$ that must be known in order to settle the outputs of matches or balance transfers. Similarly to sk^{match} , this secret key is typically controlled by a relayer.
- ▶ $pk^{\text{view}} \in \mathbb{G}$ is the public key that corresponds to a secret key $sk^{\text{view}} \in \mathbb{F}$ that must be known in order to view the contents of an encrypted wallet W .
- $r \in \mathbb{F}$ is a random secret that is used to cryptographically hide the commitments and nullifiers.

As mentioned previously, the local lists B and O allow traders to keep all relevant balance and order information private from third-parties. In Section 2.2, we explain how the randomness r prevents leakage of any information about a user's wallet.

Additionally, in Section 4.1, we show how the four public keys pk^{root} , pk^{match} , pk^{settle} , and pk^{view} allow for various levels of access control over the wallet W , and in Section 4.2, we explain the role of the fee list F .

2.2 The Commitment Tree

In order to keep track of which off-chain wallets are valid, the smart contract maintains an append-only Merkle tree of **commitments**.

Let $H : \mathbb{F} \rightarrow \mathbb{F}$ be a SNARK-friendly hash function, and let $\mathcal{H} : \mathbb{F}^n \rightarrow \mathbb{F}$ be the Merkle hash function generated by iteratively applying H . The commitment $C(W)$ to a wallet W is defined as

$$C(W) := H(\mathcal{H}(B) \parallel \mathcal{H}(O) \parallel \mathcal{H}(F) \parallel \mathcal{H}(K) \parallel r),$$

where \parallel denotes concatenation.

To perform operations on their wallet (depositing and withdrawing, submitting and cancelling orders, etc.), the user must **reveal** some nullifying information about their old wallet and **commit** to a new wallet. To ensure that the update is valid, the user must also supply a zero-knowledge proof that the update follows the rules of the protocol (e.g. does not add free tokens to B) and that the new commitment is correctly computed.

A wallet is considered valid if it has not been nullified and its commitment exists somewhere in the global Merkle tree.

Note that the randomness $r \in \mathbb{F}$ must be included in the definition of a wallet W and in the computation of the wallet commitment $C(W)$ in order to **hide** the contents of the wallet: If no such randomness were used, then adversaries could generate a rainbow table of common wallets (e.g., the wallet with zero balances and zero orders could be easily identified).

Zero-knowledge proofs are stateless (i.e., if a ZKP is valid once, it will always be valid), so the contract needs to maintain some state in order to ensure that the user cannot double-reveal a wallet that was only committed to once. To do this, when revealing an old wallet, the user computes two **nullifiers** of a wallet in addition to computing the commitment to the new wallet.

The **nullifier set** is the set of all nullifiers that have been “seen”, meaning that they have been used to reveal a wallet in the past. The contract will reject commit-reveal transactions if any of the nullifiers have been seen before.

In order to reveal their wallet W , the user first constructs a new wallet W' with the appropriate changes (a new set of orders, a change in balances to reflect an order settlement, etc.). The user then computes the two nullifiers of their old wallet W , a **wallet-spend nullifier** and a **wallet-match nullifier**.

The wallet-spend nullifier is

$$N^{\text{wallet-spend}}(W) := H(C(W) \parallel r)$$

and the wallet-match nullifier is

$$N^{\text{wallet-match}}(W) := H(C(W) \parallel r + 1).$$

We will see in Section 3.4 how this dual-nullification allows for us to perform pairwise matches between the orders in two different wallets.

Now, the user submits a zero-knowledge proof that, among other constraints, the following are true:

- There exists some valid Merkle path to $C(W)$, implying that a previous transaction committed to the wallet W .
- The nullifiers are properly computed for the wallet W .
- The transition from W to W' is valid (e.g., the user has not increased balances without depositing additional tokens).
- The user knows sk^{root} .

The contract checks that the ZK proof is valid and that the two nullifiers have not already been seen. If this check passes, the contract marks the nullifiers as seen and inserts the new commitment $C(W')$ into the commitment Merkle tree.

Illustrated in Figure 1, this basic reveal-commit scheme is used for all possible operations on a user's wallet.

In addition to the wallet commitments, the global Merkle tree also accepts insertions of **notes**. A note is essentially an unspent transaction output (i.e. a claim on some funds that




Figure 1. Commit-Reveal Scheme

needs to be redeemed). Notes can come from one of three sources:

- Settlements of matched orders between wallets.
- Internal transfers from another trader within the dark pool.
- Fees that are paid to relayers upon successful matches.

Correspondingly, when a note is redeemed, a **note-redeem** nullifier is revealed. We describe the note redemption (i.e. settlement) process further in Section 3.6.

2.3 Entire State

We have seen how *wallets* kept secret by individual traders, when combined with the idea of the *commitment tree*, allows for privately-held state with global consensus about validity of that distributed state.

In Table 1, we summarize the previous two sections into a precise description of all state (with types) held by both individual clients and the global contract.

3 Trade Lifecycle

In this section, we will go through an entire lifecycle of a trade, including creating a wallet, depositing into the system, peer discovery and handshakes in our p2p protocol, the multi-party computation itself, and settlement of matched trades.

Note that throughout this section we will use the term “trader” to refer to the party that performs these operations; however, in practice, the end-trader instead delegates much of this computation to a trustless stand-in party, termed their *relayer*. We discuss more about relayers in Section 4.

3.1 Creating a New Wallet

When a user joins Renegade for the first time, they do not have a wallet that has been committed in the global Merkle tree. So, the smart contract has a special functionality that allows for inclusion of a new wallet W without revealing any nullifiers, so long as the user proves that the wallet W is indeed a new wallet.

Specifically, the user generates a wallet

$$W = (B, O, F, K, r)$$

by setting B, O, F to be the lists of all zeros, setting K to be the tuple of appropriate access control keys as discussed in Section 4.1, and choosing a random r .

Then, this user submits $C(W)$ to the contract, along with a proof that $C(W)$ was indeed computed by committing to some wallet that had zero balances and orders. Once the contract verifies this proof, it inserts $C(W)$ into the global Merkle tree, now creating a usable wallet for the new trader.

We instantiate this argument of knowledge as a formal NP statement VALID WALLET CREATE, defined in Section A.1.

3.2 Updating a Wallet

Now that a user has committed to a new wallet W , they may *update* this wallet. When updating a wallet, a trader may do any subset of the following:

- Deposit or withdraw external ERC-20 balances from outside the dark pool.
- Send some tokens to a different user inside of the dark pool.
- Add new orders or cancel old orders in O .
- Add new fee approvals or revoke old approvals in F .

Formally, the user generates a new wallet

$$W' := (B', O', F', K', r')$$

with arbitrary O', F' , and with $r' = r + 2$. The user also creates two tuples of *transfer tokens*, the **internal transfer tuple** and the **external transfer tuple**.

The internal transfer is a tuple

$$T_I := (\tilde{m}_I, \tilde{v}_I) \in \mathbb{F}^2$$

and the external transfer is a tuple

$$T_E := (\tilde{m}_E, \tilde{v}_E, \tilde{d}_E) \in \mathbb{F}^3.$$

These two tuples determine what token (if any) to be either transferred to another user inside of the dark pool or deposited/withdrawn from the protocol entirely. \tilde{m}_I and \tilde{m}_E denote the token mint, and \tilde{v}_I and \tilde{v}_E determine the amount of tokens to be transferred. \tilde{d}_E denotes the direction (0 is deposit, 1 is withdraw) of the external transfer.

Note that either tuple may consist entirely of zeros, indicating that the user does not desire to transfer any tokens. Also, note that we cannot send T_I in-the-clear to the smart contract, or else privacy would be compromised. Instead, we encrypt T_I under the receiver’s public settle key $\text{pk}_{\text{receiver}}^{\text{settle}}$.

Then, the user submits

$$C(W'), T_E, E_{\text{pk}_{\text{receiver}}^{\text{settle}}}(T_I), N^{\text{wallet-spend}}(W), N^{\text{wallet-match}}(W)$$

to the contract, alongside a proof that W' was indeed formed by correctly applying the transfer tuples T_I and T_E to the old

	Notation	Type
CLIENT STATE		
Balances List	$B = (m_i, v_i)_{i \in [M_B]}$	\mathbb{F}^{2M_B}
Orders List	$O = (\lambda_i, m_{1_i}, m_{2_i}, s_i, p_i, a_i, \alpha_i, \tau_i)_{i \in [M_O]}$	\mathbb{F}^{8M_O}
Fees List	$F = (\text{pk}_{\text{cluster}_i}^{\text{settle}}, \mu_i, \delta_i, \gamma_i)_{i \in [M_F]}$	$(\mathbb{G} \times \mathbb{F}^3)^{M_F}$
Keys List	$K = (\text{pk}^{\text{root}}, \text{pk}^{\text{match}}, \text{pk}^{\text{settle}}, \text{pk}^{\text{view}})$	\mathbb{G}^4
Randomness	r	\mathbb{F}
CONTRACT STATE		
Merkle Path	current_merkle_path	$\mathbb{F}^L \times \mathbb{B}^L$
Nullifier Set	is_nullifier_used	$\mathbb{F} \rightarrow \mathbb{B}$
Wallet Store ⁴	wallet_store	$\mathbb{G} \rightarrow \mathbb{G}^2 \times \mathbb{F}^{2M_B+8M_O+5M_F+9}$

Table 1. Full Client and Contract State

wallet W and all commitments and nullifiers are correctly computed. As before, we provide a formal NP statement of VALID WALLET UPDATE in Section A.2.

3.3 Handshakes

Now that a user has a wallet W with non-zero lists of balances B and orders O , they may begin searching for potential counterparties to trade with.

In order to find peers, Renegade implements an off-chain **peer-to-peer messaging protocol**. Implemented over QUIC transport⁵ for high throughput, the p2p network allows for both gossip for peer discovery and message routing via a Kademlia distributed hash table.⁶

To find peers, the trader connects to the network and selects an order

$$o = (\lambda, m_1, m_2, s, p, a, \alpha, \tau) \in O$$

that they would like to match. The trader then finds the balance $b = (m, v) \in B$ that **covers** this order (e.g., if o is a buy order, then $m = m_2$). In addition, the trader selects a **relayer fee**

$$f = (\text{pk}_{\text{cluster}}^{\text{settle}}, \mu, \delta, \gamma)$$

to be taken by the party that is performing this computation.

Now, the trader generates three values $H_o = H(o \parallel r)$, $H_b = H(b \parallel r)$, and $H_f = H(f \parallel r)$. These are hiding and binding commitments to the chosen order, associated covering balance, and fee tuple; they used for cross-input consistency between the MPC and the zero-knowledge proof.

⁴The wallet store maps an ElGamal public view key (of type \mathbb{G}) to an encryption of a symmetric key (of type \mathbb{G}^2), alongside the symmetric encryption of the wallet (of type $\mathbb{F}^{2M_B+8M_O+5M_F+9}$).

⁵Langley et al., *The QUIC Transport Protocol: Design and Internet-Scale Deployment*, <https://dl.acm.org/doi/10.1145/3098822.3098842>

⁶Maymounkov and Mazières, *Kademlia: A Peer-to-peer Information System Based on the XOR Metric*, <https://www.scs.stanford.edu/dm/home/papers/kpos.pdf>

In addition, the trader computes $H_r = H(r)$, the hash of their wallet randomness, to be used in the future to cryptographically hide match outputs.

Finally, the trader generates a zero-knowledge proof π of the statement VALID COMMITMENTS, as defined in Section A.3. This statement essentially proves that the trader does indeed know some unspent wallet W containing an order o , balance b , fee f , and randomness r with the given hashes H_o , H_b , H_f , and H_r , and that there exist enough funds in the wallet W to pay for the chosen relayer fee f .

Note that the generation of π may be done completely asynchronously and be reused over multiple attempted handshakes, as the proof does not depend on the counterparty.

Now that the trader has a generated a proof π of VALID COMMITMENTS, they may begin handshaking with potential counterparties.⁷ The trader sends the handshake tuple

$$H := (\pi, N^{\text{wallet-match}}(W), \text{pk}_{\text{cluster}}^{\text{settle}}, H_o, H_b, H_f, H_r)$$

to a potential counterparty, and the counterparty checks that the proof is correct, that the wallet-match nullifier has not yet been seen on-chain (meaning that the wallet would be already spent), and that the tuple of commitments

$$(H_{o_1}, H_{b_1}, H_{f_1}, H_{o_2}, H_{b_2}, H_{f_2})$$

has not already been cached as a non-match.

If the proof is accepted by the counterparty and the order commitment pair has not already been cached, then the counterparty responds with a similar handshake tuple H_2 . The trader checks that the counterparty's handshake proof is valid, and if so, the traders proceed with the MPC.

⁷In order to bootstrap connection into the p2p network, Renegade maintains an ENS record of **authoritative relayers** that allows for new entrants to find counterparties.

3.4 MPC and Match Proofs

In order to match two party's orders, we proceed in three phases: First, perform a secret-sharing-based MPC that implements matching engine execution between the orders. Then, without opening the secret-shared output, feed this execution trace into a collaborative proof of the NP statement VALID MATCH MPC. Finally, open the outputs of both the matching engine and the collaborative proof.

This three-phase matching process allows us to only compute secret-shares of all relevant matching information, then atomically open all secret-shares at once.

Let Party 1 hold order $o_1 \in \mathbb{F}^8$, balance $b_1 \in \mathbb{F}^2$, fee $f_1 \in \mathbb{F}^5$, and randomness $r_1 \in \mathbb{F}$. Let o_2, b_2, f_2, r_2 be defined similarly for Party 2. The parties Shamir secret-share all eight of these values with each other.

Now, given all the shares⁸

$$[o_1], [b_1], [f_1], [r_1], [o_2], [b_2], [f_2], [r_2],$$

the parties run a SPDZ-style maliciously-secure MPC-with-abort to compute a secret-share of the **match tuple**

$$M := (\hat{m}_1, \hat{m}_2, \hat{v}_1, \hat{v}_2, \hat{d}, f_1, f_2) \in \mathbb{F}^{15}$$

and a secret-share of a hiding commitment to the tuple

$$H_M := H(M \parallel H_{r_1} + H_{r_2}).$$

This tuple gives the matched values between the two orders o_1 and o_2 when constrained by the balances b_1 and b_2 : That is, assuming the orders are of the same quote/base pair, \hat{v}_i is the amount of \hat{m}_i that is swapped between the two parties for $i = 1, 2$. $\hat{d} \in \mathbb{B}$ is the direction of the transfer, where $\hat{d} = 0$ means that Party 1 can increase their balances by \hat{v}_1 units of m_1 and decrease their balances by \hat{v}_2 units of m_2 , and vice-versa for Party 2.

We include the randomness $H(r_1) + H(r_2)$ in the computation of the matches commitment in order to prevent similar rainbow table-style attacks against commitments to match tuples.

Importantly, the MPC functionality that computes M must re-compute all six of the commitments

$$H_{o_1}, H_{b_1}, H_{f_1}, H_{o_2}, H_{b_2}, H_{f_2}$$

and **zero out** the output matches M in case either party lies about their inputs to the MPC. Since the hashes used in the commitment function are preimage-resistant, it is infeasible for either of the parties to manipulate their inputs without also changing their commitments.

Importantly, note that the parties *do not open* M immediately. If the parties were to open the match now, both parties

⁸In addition to the secret-shared inputs described here, the parties must also agree on a vector of **midpoint prices** from an oracle for a fixed number of assets. Note that these midpoint prices (m_i, p_i) must be revealed as public variables in the proof of VALID MATCH MPC, as the contract must assert that the prices are reasonable.

would learn information about each others' orders while being able to hangup the connection.

Now that the parties have secret-shares of every single wire in the MPC matching functionality including the output M , they perform a **collaborative proof** that produces a secret-share of a proof π_M of the statement VALID MATCH MPC. Defined formally in Section A.5, this statement essentially proves that given the parties' collective inputs o_1, b_1, o_2, b_2 , the commitment H_M is indeed the commitment to the output matches, when the matching engine is run on the input orders and balances.

Now, the traders may finally open π_M and M (alongside auxiliary MPC outputs H_M, Z_1 , and Z_2 , as defined in Section A.6), revealing the output of the matches and a proof that the matches lists were correctly computed from valid pair of orders.⁹

We illustrate this entire handshake and matching process in Figure 2.

3.5 Encumbering

Once the proof π_M , matches tuple M , and auxiliary outputs H_M, Z_1 , and Z_2 have been opened, both parties now have enough information to complete the match and the communication link between the parties may be closed.

As a final step before this proof may be submitted to the contract, the trader must prepare a few slightly altered matches lists, termed **notes**. These notes are essentially per-trader records of what tokens were swapped and are directly generated from M . One note is generated for each of the traders' wallets:

$$N_1 := ((\hat{m}_1, \hat{v}_1, \hat{d}), (\hat{m}_2, \hat{v}_2, 1 - \hat{d}), (\mu_1, \delta_1, 1), \omega, \hat{r}) \in \mathbb{F}^{11}$$

and

$$N_2 := ((\hat{m}_1, \hat{v}_1, 1 - \hat{d}), (\hat{m}_2, \hat{v}_2, \hat{d}), (\mu_2, \delta_2, 1), \omega, \hat{r} + 1) \in \mathbb{F}^{11},$$

where $\hat{r} = H_{r_1} + H_{r_2}$

Each note consists of three 3-tuples, plus two field elements ω and a randomness. Each of the three 3-tuples consist of a mint (i.e. token) address, the value of that token to be transferred in/out of the user's wallet, and the direction of the transfer.

The flag $\omega \in \mathbb{B}$ equals 1 if the note is generated from a match on a wallet, and equals 0 if the note is generated from either a fee cut or an internal transfer. Here, since these notes N_1 and N_2 are the notes from a wallet match, we set $\omega = 1$.

⁹The naive way to open the MPC is for each party to send their secret-shares to each other. This scheme fails to have **opening fairness**, however, as the first party to receive the other's shares can simply hang up the connection without replying with their corresponding shares. To remedy this, we optionally allow for opening to be routed via an arbitrary trusted third party to ensure fair opening.

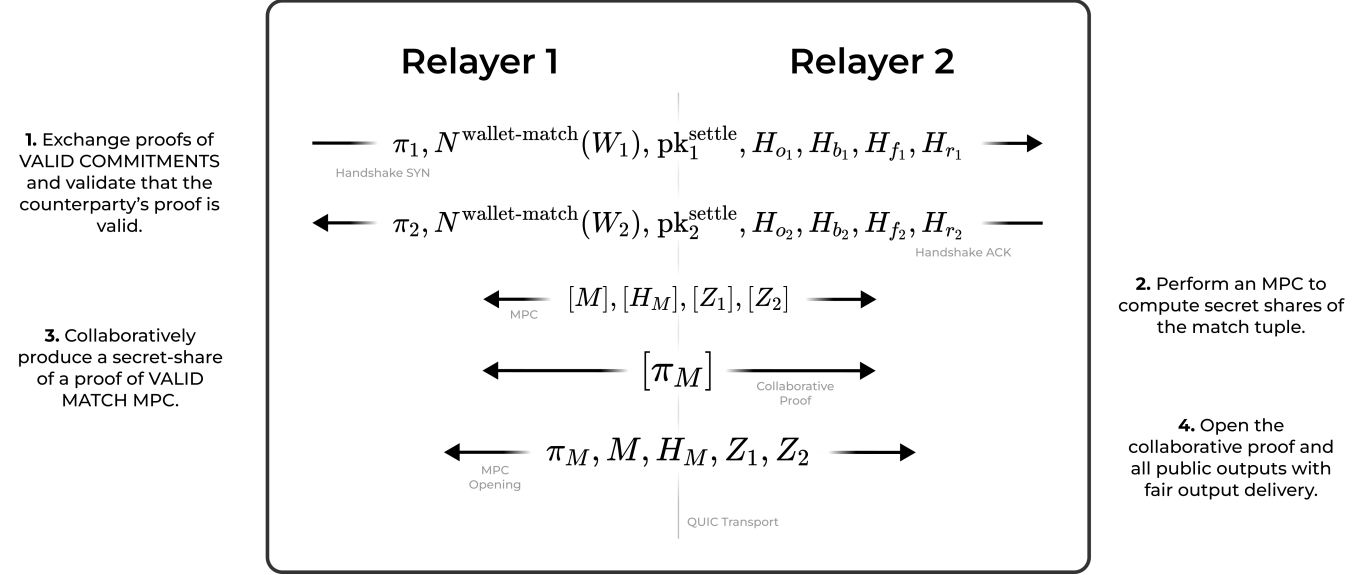


Figure 2. Inter-Relayer Communication Flow

In addition to these two notes, we generate two extra notes for the relayer fees, plus one final note for the global protocol fee.

We cannot send notes in-the-clear to the contract (or else third-parties would learn what tokens are being swapped), so we encrypt¹⁰ these two matches lists under the corresponding settle key of each party:

$$E_{\text{pk}_1^{\text{settle}}}(N_1) \text{ and } E_{\text{pk}_2^{\text{settle}}}(N_2)$$

Now, the trader generates a proof π_E of the statement VALID MATCH ENCRYPTION as defined in Section A.7, which essentially shows that the notes N_1, N_2 we indeed formed from the original match tuple M , that M is indeed consistent with the publicly-known commitment H_M , and that the notes are properly encrypted.

Finally, after generating π_E , the trader is now ready to interact with the smart contract. The trader sends four different proofs $\pi_1, \pi_2, \pi_M, \pi_E$ to the contract (i.e. two proofs of VALID COMMITMENTS, one proof of VALID MATCH MPC, and one proof of VALID MATCH ENCRYPTION), alongside the union of public variables for all four proofs.

The contract checks that all four zero-knowledge proofs are valid under the given public inputs. If all checks pass, the contract then marks the two nullifiers $N^{\text{wallet-match}}(W_1)$ and $N^{\text{wallet-match}}(W_2)$ as being seen, and inserts all of the encrypted notes into the commitment tree.

Note that this nullification is different from how reveal-commit schemes work in VALID WALLET UPDATE as in

¹⁰In our instantiation of the protocol, we use ElGamal to encrypt all notes. In addition to being a SNARK-friendly encryption scheme, ElGamal has the property of being **key-private**, meaning that third-parties cannot even see which public settle key of the party who is receiving the funds.

Section 3.2. Here, we only mark the *wallet-match* nullifiers as seen, not the *wallet-spend* nullifiers. In addition, we do not insert a commitment to a wallet into the global Merkle tree; rather, we are inserting encrypted notes.

Note that since we have revealed the wallet-match nullifiers, calling VALID WALLET UPDATE is now impossible, as neither party can prove that their wallet has an unseen wallet-match nullifier. Both wallets are now considered “encumbered”, and the only operation that either party may perform is to settle their matched order.

3.6 Settlement

Now that the trader’s wallet W has been matched and encumbered, they need to **settle** this match in order to update their balances and un-encumber the wallet.

To do this, the trader first obtains a note

$$N = \left((\hat{m}_1, \hat{v}_1, \hat{d}_1), (\hat{m}_2, \hat{v}_2, \hat{d}_2), (\hat{m}_3, \hat{v}_3, \hat{d}_3), \omega, \hat{r} \right).$$

The trader can find the note by either remembering the match they just performed, or if the note was generated from an internal transfer from a different user inside the pool, by scanning through the commitment history to find the most recent encrypted note and decrypt it under their secret key $\text{sk}^{\text{settle}}$.

Now, the trader constructs a new wallet

$$W' = (B', O', F', K', r')$$

such that F' and K' are unchanged from the original wallet W , and with $r' = r + 2$.

To construct B' , the trader simply adds or subtracts values v_i from the balances list according to the note N . If the trader

is settling a matched order, there will always be exactly one balance that is increased, and exactly one balance that is decreased.

Finally, to construct O' , the trader finds the order

$$o = (\lambda_i, m_{1,i}, m_{2,i}, s_i, p_i, a_i, \alpha_i, \tau_i) \in O$$

that was matched by N and decreases the size a_i by the corresponding matched value \hat{v}_1 or \hat{v}_2 depending on the direction of the match.

Now, given this new wallet W' that was formed by directly settling the match note N against the old wallet W , the trader constructs a proof π_S of the statement VALID SETTLE as defined in Section A.8.

In addition to revealing the commitment to the new wallet $C(W')$ and the wallet-spend and wallet-match nullifiers as normal, the trader also reveals a **note-redeem** nullifier defined as

$$N^{\text{note-redeem}}(N) := H(N \parallel \text{pk}^{\text{settle}}).$$

Note-redeem nullifiers exist in order to prevent replay-style attacks by double-settling a matched order.

The trader then sends this proof π_S to the contract, and assuming it is correctly verified, the contract will mark both the wallet-spend and note-redeem nullifiers as being seen, and insert the new commitment $C(W')$ into the Merkle tree. Note that if $\omega = 0$ (i.e., the note came from an internal transfer or fee output), then the contract also asserts that the wallet-match nullifier has not been seen, and marks it as seen. This allows the system to have a single operation (settlement) for all notes, no matter how the note was generated.

Now, the trader has settled their matched orders, and all three nullifiers (wallet-match, wallet-spend, note-redeem) have been seen, making further use of ether the old wallet W or the used note N impossible.

In summary, this basic update-match-settle lifecycle is used for every single order that a trader would like to perform, all while avoiding any information leakage to third-parties.

4 Relay Delegation

In the previous Section 3 outlining the entire lifecycle of a trade, we assumed that the end-trader would be online at all times to handshake and perform MPC calculations with arbitrary counterparties.

However, this is unreasonably restrictive: Traders may want to “fire and forget” their orders, much like how current centralized exchanges operate. In addition, the trader may want to handshake with many counterparties at once (every single other node in the network may have a valid counter-order).

To allow for this, we introduce the concept of **relayers**. A relayer is essentially a stand-in for a trader that holds




Figure 3. Key Derivation Hierarchy

their wallet W and continually attempts MPC calculations followed by match-settle proofs in the event of a match.

In Section 4.1, we describe the key hierarchy that allows for trust minimization between the end-trader and their relayer, and in Section 4.2 we describe the high-level network topology of the system, including the idea of “relayer clusters”. Finally, in Section 4.3, we describe the role of the fee list F in compensating the relayer.

4.1 Key Hierarchy

In designing the relayer system, the primary goal is *trust minimization* between the end-trader and the relayer. Specifically, a relayer should only ever be able to match outstanding orders and settle previous matches; the relayer should *not* be able to create or cancel orders, or deposit or withdraw funds.

To implement this level of access control, we introduce the **key hierarchy**, as outlined in Figure 3. Alongside the base Ethereum keypair, we have four different levels of Renegade-native keys, all with various levels of access controls.

The key hierarchy begins with a trader’s Ethereum keypair $(\text{pk}^{\text{eth}}, \text{sk}^{\text{eth}})$ on the secp256k1 curve. Let

$$(r, s, v) \in \mathbb{B}^{256} \times \mathbb{B}^{256} \times \mathbb{B}^8$$

be the deterministic signature of the message `create_sk_root`. Then, construct

$$\text{sk}^{\text{root}} = H(r \parallel s \parallel v)$$

and recover the corresponding public key pk^{root} .

Using this process, we may always re-derive the root keypair so long as the trader does not lose their native Ethereum keypair. Note that we may skip this derivation entirely and simply generate a random sk^{root} if the trader does not want to maintain an Ethereum keypair, or if the trader wants to use a Renegade-native multisig solution.

This **root keypair** is the ultimate authority over a user’s wallet: Any user who knows this secret key may perform

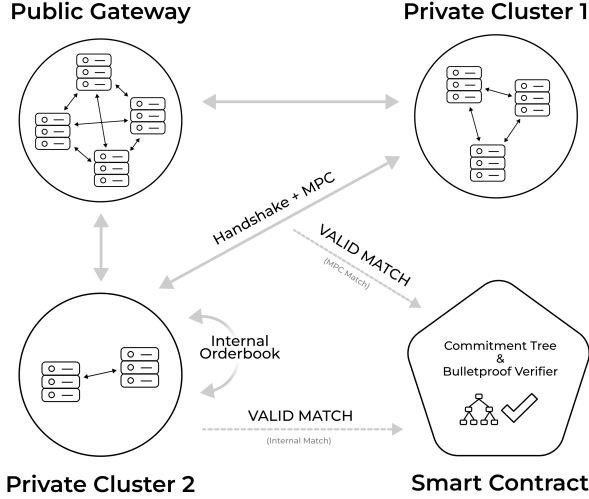


Figure 4. Network Topology, showing matches derived from both MPC and internal proofs.

arbitrary operations to the balances and orders inside the wallet, including withdrawing all funds.

Now that we have sk^{root} , we derive the other three key-pairs in a very similar way:

$$\begin{aligned} sk^{match} &= H(\text{sign}_{sk^{root}}(\text{create_sk_match})) \\ sk^{settle} &= H(\text{sign}_{sk^{match}}(\text{create_sk_settle})) \\ sk^{view} &= H(\text{sign}_{sk^{settle}}(\text{create_sk_view})) \end{aligned}$$

and correspondingly recover the public keys pk^{match} , pk^{settle} , and pk^{view} . These three final keys are the **match keypair**, the **settle keypair**, and the **view keypair**.

Any party who knows the match keypair is allowed to match outstanding orders in the orders list O ; importantly, this key is *not* able to arbitrarily modify O , deposit, or withdraw. The settle keypair has less authority, only being able to settle previous matches, or settle direct token transfers from another user inside the pool. Finally, the view keypair has the least amount of authority, only being able to decrypt and view the wallet; the view keypair cannot modify the wallet in any way.

By introducing this five-level key hierarchy, the system allows for various levels of access control. In particular, the trader *only* sends sk^{match} to the relay who will match orders on the trader’s behalf. Even if the relay is completely compromised, the worst outcome is that the relay leaks the wallet W , compromising the privacy of the trader, but not stealing funds.

4.2 Network Architecture

In Figure 4, we illustrate the high-level p2p network topology.

At the base layer, the network simply consists of various relayers that communicate with each other. The network is permissionless, meaning that at any time any new trader may enter the network as their own relayer and begin MPC handshakes.

However, since the network needs to support a large number of potential matches on many orders, the relayers are grouped into logical units called **relayer clusters**.

These clusters are fault-tolerant replicated groups of relayer nodes that all manage the same set of wallets. This replication allows for higher throughput of matches (since many relayers can try different matches at once), and allows for fault-tolerance under network interruption and partition.

In Figure 4, we illustrate three different clusters communicating with each other: PRIVATE CLUSTER 1 with three relayers, a smaller PRIVATE CLUSTER 2 with two relayers, and a special larger PUBLIC GATEWAY with four relayers. The Public Gateway is a cluster of relayers like the rest (i.e., it has no special permissions), but is a Renegade-operated set of relayers that allows for bootstrap connectivity into the network, and allows for traders who do not want to run their own infrastructure to participate in the network.

Note that privacy-concerned high-volume traders should run their own relay clusters, as using the Gateway exposes trader’s wallets to the centralized Gateway provider.

In addition to the standard handshake and MPC process, Figure 4 also illustrates the idea of an “internal match”, to be described in Section 5. Also, Figure 4 illustrates the submission of a VALID MATCH MPC proof to the global Merkle commitment tree and zero-knowledge-proof verification contract.

4.3 Relayer Fees

In the definition of a wallet W , there is one final element that we have not yet described: The **fee list** F . Since running a relayer requires somewhat expensive hardware (e.g., zero-knowledge proving is quite memory-intensive), and since smart contract calls require some protocol fee, relayers naturally need compensation for performing matching and settlement.

To implement this functionality, each relayer that wants to match public wallets advertises a tuple

$$(pk_{cluster_i}^{settle}, \mu_i, \delta_i, \gamma_i).$$

The key $pk_{cluster_i}^{settle}$ is the settle key of some wallet controlled by a relay cluster. μ_i is the mint address of the token that constant fees are denominated in. This is typically either some stablecoin or some L2-native token to pay for fees. The first fee parameter δ_i is the constant fee that the relayer may take upon any successful operation on the user’s wallet (matching, settling, etc.), and the second fee parameter γ_i is the percentage fee that they relayer may take on each match.

If this key-fee tuple is acceptable, the trader inserts this tuple into their fee list F . We include a formalization of this fee-taking process in Section A.8.

Importantly, note that fees may be avoided entirely for traders who run their own relay clusters: Indeed, we charge sizeable fees on the Public Gateway to promote maximal decentralization of the network.

5 Indications of Interest

As mentioned in Section 2, the principal goal of the base-layer Renegade system is to ensure complete privacy of all relevant values (orders, balances, matches, etc.) from all third-parties to the trade.

However, in practice, having completely obfuscated “dark” orders may not be optimal for liquidity provision: Indeed, there is a core tradeoff between quality-of-execution and speed-of-execution (dark pools give best price, whereas lit pools let you transact immediately). For traders who may want to tradeoff price for speed, Renegade allows for additional **indications of interest** flags.

An indication of interest is some predicate on a wallet W , proved in zero-knowledge. For example, a trader may reveal the fact “my wallet W is a buy order of WETH/USDC at the midpoint price”, without disclosing the size of the trade. To ensure that the trader is not lying, this predicate must be proved in zero-knowledge as a part of the handshake process.

In Section A.4, we give formal NP statments for every IoI flag that Renegade supports, including the base/quote token pair, the side (buy or sell), if the order is a midpoint-pegged order, the limit price of the order, and the size of the order.

Note that the maximal indication of interest (i.e., turning on and proving all IoI flags) makes an order *equivalent to a lit order*. Indeed, we can embed an entire lit orderbook with dark-lit crossover matches inside of Renegade.

Note that if a counterparty turns on all IoI flags, there is actually no need to compute a MPC as normal: We may directly lift this lit order and match it with one of our own orders. In Section A.6, we formally define the NP statement VALID MATCH LIT that allows for this functionality.

Finally, one optimization to note is that if a relayer manages two different wallets W_1 and W_2 that contain an overlapping order, once again MPC is not necessary and the relayer may simply match these two orders directly. Here, the relayer computes the same VALID MATCH MPC as normal, but does not need to actually run a MPC with itself in order to generate the proof π_M .

6 Conclusion

Renegade aims to solve the four core problems outlined in Section 1: MEV, pre-trade transparency, post-trade transparency, and address discrimination. We claim the following two strong privacy properties as solutions to these problems:

- All third parties who are neither the end-user nor a managing relayer learn zero information about the activities inside of the pool, other than global token inflows and outflows. In particular, third parties cannot deduce the balances or orders of any wallet, and they cannot deduce any details of any match or settlement other than the fact that a match and settlement occurred between some pair of traders.
- Individual relayers learn nothing about the state of balances and orders of wallets that they do not manage. The Public Gateway is not a special relay cluster and has no in-protocol advantages over private clusters.

In all, we have seen how we can combine the ideas of a commit-reveal scheme, pairwise MPC, and zero-knowledge settlement to create a protocol that is functionally equivalent to a CLOB, yet maximally private; RENEGADE allows for a truly anonymous global exchange of value.

A Formal NP Statements

Here, we give precise specifications for each of the eight different NP statements that are used by the protocol. All of these statements are implemented as rank-one constraint systems and proved/verified via Bulletproofs.¹¹

Each statement consists of a list of “public variables” that are publicly known to the verifier. The “private variables” are the secret witnesses known only by the prover. Finally, we have a “such that” list of all properties that are encoded into the constraint systems.

Note: For notational clarity, we have omitted the range constraints. These are added by simply constraining all input variables (both public and private) to be less than 2^{128} .

A.1 VALID WALLET CREATE

WITH PUBLIC VARIABLES

- $C(W') \in \mathbb{F}$, the commitment to the new wallet.
- $E_{\text{pk}^{\text{view}}}(W') \in \mathbb{F}^{2M_B+8M_O+5M_F+9}$, the encryption of W' under pk^{view}

I KNOW PRIVATE VARIABLES

- $F \in \mathbb{F}^{4M_F}$
- $K = (\text{pk}^{\text{root}}, \text{pk}^{\text{match}}, \text{pk}^{\text{settle}}, \text{pk}^{\text{view}}) \in \mathbb{G}^4$
- $\text{sk}^{\text{root}}, \text{sk}^{\text{match}}, \text{sk}^{\text{settle}}, \text{sk}^{\text{view}} \in \mathbb{F}$
- $r \in \mathbb{F}$

SUCH THAT

- $C(W') = C(0^{2M_B}, 0^{8M_O}, F, K, r)$
- $E_{\text{pk}^{\text{view}}}(W')$ is the proper ElGamal encryption of W' under the public key pk^{view}
- pk^{root} is the valid public key corresponding to sk^{root}
- pk^{match} is the valid public key corresponding to sk^{match}
- $\text{pk}^{\text{settle}}$ is the valid public key corresponding to $\text{sk}^{\text{settle}}$
- pk^{view} is the valid public key corresponding to sk^{view}
- $\text{sk}^{\text{match}} = H(\text{sign}_{\text{sk}^{\text{root}}}(\text{create_sk_match}))$
- $\text{sk}^{\text{settle}} = H(\text{sign}_{\text{sk}^{\text{match}}}(\text{create_sk_settle}))$
- $\text{sk}^{\text{view}} = H(\text{sign}_{\text{sk}^{\text{settle}}}(\text{create_sk_view}))$

¹¹Bünz et al., *Bulletproofs: Short Proofs for Confidential Transactions and More*, <https://eprint.iacr.org/2017/1066>

A.2 VALID WALLET UPDATE

WITH PUBLIC VARIABLES

- $C(W') \in \mathbb{F}$, the commitment to the new wallet.
- $E_{\text{pk}^{\text{view}}}(W') \in \mathbb{G}^2 \times \mathbb{F}^{2M_B+8M_O+5M_F+9}$, the encryption of W' under pk^{view}
- $N^{\text{wallet-spend}}(W) \in \mathbb{F}$
- $N^{\text{wallet-match}}(W) \in \mathbb{F}$
- $R_{\text{global}} \in \mathbb{F}$, the current root of the commitment Merkle tree.
- $\text{pk}_{\text{receiver}}^{\text{settle}}$, the settle public key of the user to receive the internal transfer.
- $E_{\text{pk}_{\text{receiver}}^{\text{settle}}}(T_I)$, the encryption of the internal transfer tuple.
- $T_E = (\tilde{m}_E, \tilde{v}_E, \tilde{d}_E) \in \mathbb{F}^3$, the external transfer tuple.
- $\tau \in \mathbb{F}$, the timestamp of this update.

I KNOW PRIVATE VARIABLES

- $W = (B, O, F, K, r)$, the old wallet.
- $W' = (B', O', F', K', r')$, the new wallet.
- $\text{Open}(C(W), R_{\text{global}})$, a Merkle proof that $C(W)$ is inserted into the commitment tree.
- $\text{sk}^{\text{root}} \in \mathbb{F}$
- $T_I = (\tilde{m}_I, \tilde{v}_I) \in \mathbb{F}^2$, the internal transfer tuple.

SUCH THAT

- $C(W)$ is correctly computed.
- $C(W')$ is correctly computed.
- $E_{\text{pk}^{\text{view}}}(W')$ is correctly computed.
- $E_{\text{pk}_{\text{receiver}}^{\text{settle}}}(T_I)$ is correctly computed.
- $N^{\text{wallet-spend}}(W)$ is correctly computed.
- $N^{\text{wallet-match}}(W)$ is correctly computed.
- $\text{Open}(C(W), R_{\text{global}})$ is a valid Merkle proof.
- $K' = K$
- $r' = r + 2$
- pk^{root} is the valid public key corresponding to sk^{root}
- $\tilde{d}_E \in \{0, 1\}$
- For all balances $b'_i = (m'_i, v'_i) \in B'$:
 - Either $m'_i = 0$, or m'_i is unique in the list of all mints of B' . (i.e., no duplicate mints are allowed).
 - v'_i is equal to

$$\sum_{j \in [M_B] \text{ s.t. } m_j = m'_i} v_j$$

plus

$$\mathbf{1}_{m'_i = \tilde{m}_E \wedge \tilde{d}_E = 0} \tilde{v}_E - \mathbf{1}_{m'_i = \tilde{m}_E \wedge \tilde{d}_E = 1} \tilde{v}_E$$

minus

$$\mathbf{1}_{m'_i = \tilde{m}_I} \tilde{v}_I$$

(i.e., balances are unchanged, except for a deposit or withdraw according to T_E and a transfer according to T_I)

- For all orders $o'_i = (\lambda'_i, m'_{1_i}, m'_{2_i}, s'_i, p'_i, a'_i, \alpha'_i, \tau'_i) \in O'$:
 - Either $o'_i = \mathbf{0}^8$, or (m'_{1_i}, m'_{2_i}) is unique in the list of all mint pairs of B' . (i.e., no duplicate token pairs are allowed).
 - If $\lambda'_i = \lambda_i$ and $m'_{1_i} = m_{1_i}$ and $m'_{2_i} = m_{2_i}$ and $s'_i = s_i$ and $p'_i = p_i$ and $a'_i = a_i$ and $\alpha'_i = \alpha_i$, then $\tau'_i = \tau_i$. Otherwise, $\tau'_i = \tau$. (i.e., if any order has changed, then the timestamp is set to the current time)

A.3 VALID COMMITMENTS

WITH PUBLIC VARIABLES

- $N^{\text{wallet-match}}(W) \in \mathbb{F}$
- $R_{\text{global}} \in \mathbb{F}$, the current root of the commitment Merkle tree.
- $\text{pk}^{\text{settle}} \in \mathbb{G}$
- $H_o, H_b, H_f \in \mathbb{F}$, hiding commitments to an order, balance, and fee tuple, respectively.
- $H_r \in \mathbb{F}$, the hash of the randomness.

I KNOW PRIVATE VARIABLES

- $C(W) \in \mathbb{F}$, the commitment to the wallet.
- $R_B, R_O, R_F \in \mathbb{F}$, Merkle roots of the balances, orders, and fees of the wallet.
- $K \in \mathbb{G}^4$, the key list of the wallet.
- $r \in \mathbb{F}$, the randomness of the wallet.
- $\text{Open}(C(W), R_{\text{global}})$, a Merkle proof that $C(W)$ is inserted into the global commitment tree.
- $\text{Open}(o, R_O)$, a Merkle proof that o is inserted into the internal Merkle tree of orders.
- $\text{Open}(b, R_B)$, a Merkle proof that b is inserted into the internal Merkle tree of orders.
- $\text{Open}(f, R_F)$, a Merkle proof that f is inserted into the internal Merkle tree of orders.
- $\text{sk}^{\text{root}} \in \mathbb{F}$
- $\text{sk}^{\text{match}} \in \mathbb{F}$
- $o = (\lambda, m_1, m_2, s, p, a, \alpha, \tau) \in O$
- $b = (m, v) \in B$
- $b' = (m', v') \in B$, the balance corresponding to the constant fee δ .
- $f = (\text{pk}_{\text{cluster}}^{\text{settle}}, \mu, \delta, \gamma) \in F$

SUCH THAT

- $N^{\text{wallet-match}}(W)$ is correctly computed.
- $C(W) = H(R_B || R_O || R_F || \mathcal{H}(K) || r)$
- $\text{Open}(C(W), R_{\text{global}})$ is a valid Merkle proof.
- Either $\text{Open}(o, R_O)$ is a valid Merkle proof, or pk^{root} is the valid public key corresponding to sk^{root} .
- $\text{Open}(b, R_B)$ is a valid Merkle proof.
- $\text{Open}(f, R_F)$ is a valid Merkle proof.
- pk^{match} is the valid public key corresponding to sk^{match}
- $H_o = H(o || r)$
- $H_b = H(b || r)$
- $H_f = H(f || r)$
- $H_r = H(r)$
- $m = s \cdot m_1 + (1 - s) \cdot m_2$
- $m' = \mu$ and $v' \geq \delta$

A.4 Indications of Interest Statements

In this section, we provide the formal specifications for a few different indications of interest. **Note:** We only provide partial statements here; each statement should be conjoined with a proof of VALID COMMITMENTS.

A.4.1 VALID IOI TYPE. Simply add “ $k \in \mathbb{F}$ ” to the public variables.

A.4.2 VALID IOI PAIR. Simply add “ $m_1, m_2 \in \mathbb{F}$ ” to the public variables.

A.4.3 VALID IOI SIDE. Simply add “ $s \in \mathbb{F}$ ” to the public variables.

A.4.4 VALID IOI LIMIT PRICE. Simply add “ $p \in \mathbb{F}$ ” to the public variables.

A.4.5 VALID IOI AMOUNT. Simply add “ $a \in \mathbb{F}$ ” to the public variables.

A.4.6 VALID IOI BALANCE BOUND. Add “ $v \in \mathbb{F}$ ” to the public variables and “ $v \geq \nu$ ” to the constraint system.

A.4.7 VALID IOI FEE. Simply add “ $f \in \mathbb{F}^5$ ” to the public variables.

A.5 VALID MATCH MPC

WITH PUBLIC VARIABLES

- $(m_i, p_i)_{i \in [M_P]} \in \mathbb{F}^{2M_P}$, the vector of midpoint oracle prices.
- $H_{o_1}, H_{b_1}, H_{f_1}, H_{r_1} \in \mathbb{F}$, the commitments to the order, balance, fee, and randomness from Relayer 1.
- $H_{o_2}, H_{b_2}, H_{f_2}, H_{r_2} \in \mathbb{F}$, the commitments to the order, balance, fee, and randomness from Relayer 2.
- $H_M \in \mathbb{F}$, the hiding commitment to the matches tuple M .
- $Z_1 \in \mathbb{F}$, the bit that equals 1 iff M is a non-trivial matches list.
- $Z_2 \in \mathbb{F}$, the **zeroing bit** that equals 1 if any party lies about their secret input w.r.t. to the public commitments $H_{o_1}, H_{o_2}, H_{b_1}, H_{b_2}, H_{f_1}, H_{f_2}$, and equals 1 otherwise.

I KNOW PRIVATE VARIABLES

- $o_1 = (\lambda_1, m_{1_1}, m_{2_1}, s_1, p_1, a_1, \alpha_1, \tau_1) \in \mathbb{F}^8$
- $o_2 = (\lambda_2, m_{1_2}, m_{2_2}, s_2, p_2, a_2, \alpha_2, \tau_2) \in \mathbb{F}^8$
- $b_1 = (m_1, v_1) \in \mathbb{F}^2$
- $b_2 = (m_2, v_2) \in \mathbb{F}^2$
- $f_1 = (\text{pk}_{\text{cluster}_1}^{\text{settle}}, \mu_1, \delta_1, \gamma_1) \in \mathbb{F}^5$
- $f_2 = (\text{pk}_{\text{cluster}_2}^{\text{settle}}, \mu_2, \delta_2, \gamma_2) \in \mathbb{F}^5$
- $M = (\hat{m}_1, \hat{m}_2, \hat{v}_1, \hat{v}_2, \hat{d}, f_1, f_2) \in \mathbb{F}^{15}$, the match tuple.

SUCH THAT

- $H_M = H(M \parallel H_{r_1} + H_{r_2})$
- $Z_1 \in \{0, 1\}$
- $Z_2 \in \{0, 1\}$
- (\hat{v}_1, \hat{v}_2) is the output of the matching engine operation on the orders o_1, o_2 under the balance constraints b_1, b_2 .
- $\hat{d} = s_1 \cdot (1 - s_2)$
- $Z_1 = (1 - Z_2) \cdot \mathbf{1}_{\hat{v}_1 \neq 0 \vee \hat{v}_2 \neq 0}$
 $Z_2 = 1 - \mathbf{1}_{H_{o_1}=H(o_1||r_1)} \cdot \mathbf{1}_{H_{o_2}=H(o_2||r_2)}$
- $\cdot \mathbf{1}_{H_{b_1}=H(b_1||r_1)} \cdot \mathbf{1}_{H_{b_2}=H(b_2||r_2)}$
 $\cdot \mathbf{1}_{H_{f_1}=H(f_1||r_1)} \cdot \mathbf{1}_{H_{f_2}=H(f_2||r_2)}$

A.6 VALID MATCH LIT

WITH PUBLIC VARIABLES

- $(m_i, p_i)_{i \in [M_P]} \in \mathbb{F}^{2M_P}$, the vector of midpoint oracle prices.
- $H_{o_1}, H_{b_1}, H_{f_1}, H_{r_1} \in \mathbb{F}$, the commitments to the order, balance, fee, and randomness from Relayer 1.
- $H_{o_2}, H_{b_2}, H_{f_2}, H_{r_2} \in \mathbb{F}$, the commitments to the order, balance, fee, and randomness from Relayer 2.
- $H_M \in \mathbb{F}$, the hiding commitment to the matches tuple M .
- $Z_1 \in \mathbb{F}$, the bit that is 1 iff M is a non-trivial matches list.
- $Z_2 \in \mathbb{F}$, the zeroing bit that equals 0 if Party 2 (the non-lit party) lies about their secret input w.r.t. to the public commitments H_{o_2}, H_{b_2} , and equals 1 otherwise.¹²
- $o_1 = (\lambda_1, m_{1_1}, m_{2_1}, s_1, p_1, a_1, \alpha_1, \tau_1) \in \mathbb{F}^8$
- $b_1 = (m_1, v_1) \in \mathbb{F}^2$
- $f_1 = (\text{pk}_{\text{cluster}_1}^{\text{settle}}, \mu_1, \delta_1, \gamma_1) \in \mathbb{F}^5$

I KNOW PRIVATE VARIABLES

- $o_2 = (\lambda_2, m_{1_2}, m_{2_2}, s_2, p_2, a_2, \alpha_2, \tau_2) \in \mathbb{F}^8$
- $b_2 = (m_2, v_2) \in \mathbb{F}^2$
- $f_2 = (\text{pk}_{\text{cluster}_2}^{\text{settle}}, \mu_2, \delta_2, \gamma_2) \in \mathbb{F}^5$
- $r_2 \in \mathbb{F}$
- $M = (\hat{m}_1, \hat{m}_2, \hat{v}_1, \hat{v}_2, \hat{d}, f_1, f_2) \in \mathbb{F}^{15}$, the match tuple.

SUCH THAT

- $H_M = H(M \parallel H_{r_1} + H_{r_2})$
- $Z_1 \in \{0, 1\}$
- $Z_2 \in \{0, 1\}$
- (\hat{v}_1, \hat{v}_2) is the output of the matching engine operation on the orders o_1, o_2 under the balance constraints b_1, b_2 , potentially after the constant fees δ_1, δ_2 have been deducted from the balance tuples.
- $\hat{d} = s_1 \cdot (1 - s_2)$
- $Z_1 = (1 - Z_2) \cdot \mathbf{1}_{\hat{v}_1 \neq 0 \vee \hat{v}_2 \neq 0}$
 $Z_2 = 1 - \mathbf{1}_{H_{o_1}=H(o_1 \parallel r_1)} \cdot \mathbf{1}_{H_{o_2}=H(o_2 \parallel r_2)}$
- $\cdot \mathbf{1}_{H_{b_1}=H(b_1 \parallel r_1)} \cdot \mathbf{1}_{H_{b_2}=H(b_2 \parallel r_2)}$
 $\cdot \mathbf{1}_{H_{f_1}=H(f_1 \parallel r_1)} \cdot \mathbf{1}_{H_{f_2}=H(f_2 \parallel r_2)}$

¹²Note that it is impossible for Party 1 (the lit party) to lie about their input, as they have proven every single IoI, and therefore we know their order o_1 and balance bound b_1 exactly.

A.7 VALID MATCH ENCRYPTION

WITH PUBLIC VARIABLES

- $H_M \in \mathbb{F}$, the hiding commitment to the matches tuple M .
- $H_{r_1}, H_{r_2} \in \mathbb{F}$, the commitments to the individual parties' randomnesss.
- $\text{pk}_1^{\text{settle}} \in \mathbb{G}$, the settle key of Party 1's wallet.
- $\text{pk}_2^{\text{settle}} \in \mathbb{G}$, the settle key of Party 2's wallet.
- $\text{pk}_{\text{protocol}}^{\text{settle}} \in \mathbb{G}$, the settle key of the global protocol fee wallet.
- $\gamma_{\text{protocol}} \in \mathbb{F}$, the global protocol fee value.
- $E_{\text{pk}_1^{\text{settle}}}(N_1)$, the encrypted note for Party 1's settlement.
- $E_{\text{pk}_2^{\text{settle}}}(N_2)$, the encrypted note for Party 2's settlement.
- $E_{\text{pk}_{\text{cluster}_1}^{\text{settle}}}(N_{R_1})$, the encrypted note for Relayer 1's fees.
- $E_{\text{pk}_{\text{cluster}_2}^{\text{settle}}}(N_{R_2})$, the encrypted note for Relayer 2's fees.
- $E_{\text{pk}_{\text{protocol}}^{\text{settle}}}(N_P)$, the encrypted note for the in-protocol fee.

I KNOW PRIVATE VARIABLES

- $M = \begin{pmatrix} \hat{m}_1, \hat{m}_2, \hat{v}_1, \hat{v}_2, \hat{d}, (\text{pk}_{\text{cluster}_1}^{\text{settle}}, \mu_1, \delta_1, \gamma_1) \\ (\text{pk}_{\text{cluster}_2}^{\text{settle}}, \mu_2, \delta_2, \gamma_2) \end{pmatrix}$,
the match tuple.
- $N_1 = \left((\hat{m}_1, \hat{v}_1^1, \hat{d}), (\hat{m}_2, \hat{v}_2^1, 1 - \hat{d}), (\mu_1, \delta_1, 1), 1, \hat{r} \right)$
- $N_2 = \left((\hat{m}_1, \hat{v}_1^2, 1 - \hat{d}), (\hat{m}_2, \hat{v}_2^2, \hat{d}), (\mu_2, \delta_2, 1), 1, \hat{r} + 1 \right)$
- $N_{R_1} = \left((\hat{m}_1, \hat{v}_1^{R_1}, 0), (\hat{m}_2, \hat{v}_2^{R_1}, 0), (\mu_1, \delta_1, 0), 0, \hat{r} + 2 \right)$
- $N_{R_2} = \left((\hat{m}_1, \hat{v}_1^{R_2}, 0), (\hat{m}_2, \hat{v}_2^{R_2}, 0), (\mu_2, \delta_2, 0), 0, \hat{r} + 3 \right)$
- $N_P = \left((\hat{m}_1, \hat{v}_1^P, 0), (\hat{m}_2, \hat{v}_2^P, 0), (0, 0, 0), 0, \hat{r} + 4 \right)$

SUCH THAT

- $H_M = H(M \parallel H_{r_1} + H_{r_2})$
- $\hat{r} = H_{r_1} + H_{r_2}$
- N_1 is properly encrypted under $\text{pk}_1^{\text{settle}}$
- N_2 is properly encrypted under $\text{pk}_2^{\text{settle}}$
- N_{R_1} is properly encrypted under $\text{pk}_{\text{cluster}_1}^{\text{settle}}$
- N_{R_2} is properly encrypted under $\text{pk}_{\text{cluster}_2}^{\text{settle}}$
- N_P is properly encrypted under $\text{pk}_{\text{protocol}}^{\text{settle}}$
- If $d = 0$, then:
 - $v_1^1 = \hat{v}_1 \cdot (1 - \gamma_1 - \gamma_{\text{protocol}})$ and $v_2^1 = \hat{v}_2$
 - $v_1^2 = \hat{v}_1$ and $v_2^2 = \hat{v}_2 \cdot (1 - \gamma_2 - \gamma_{\text{protocol}})$
 - $v_1^{R_1} = \hat{v}_1 \cdot \gamma_1$ and $v_2^{R_1} = 0$
 - $v_1^{R_2} = 0$ and $v_2^{R_2} = \hat{v}_2 \cdot \gamma_2$
- If $d = 1$, then:
 - $v_1^1 = \hat{v}_1$ and $v_2^1 = \hat{v}_2 \cdot (1 - \gamma_1 - \gamma_{\text{protocol}})$
 - $v_1^2 = \hat{v}_1 \cdot (1 - \gamma_2 - \gamma_{\text{protocol}})$ and $v_2^2 = \hat{v}_2$
 - $v_1^{R_1} = 0$ and $v_2^{R_1} = \hat{v}_2 \cdot \gamma_1$
 - $v_1^{R_2} = \hat{v}_1 \cdot \gamma_2$ and $v_2^{R_2} = 0$
- $v_1^P = \hat{v}_1 \cdot \gamma_{\text{protocol}}$
- $v_2^P = \hat{v}_2 \cdot \gamma_{\text{protocol}}$

A.8 VALID SETTLE

WITH PUBLIC VARIABLES

- $C(W') \in \mathbb{F}$, the commitment to the new wallet.
- $E_{\text{pk}^{\text{view}}}(W') \in \mathbb{G}^2 \times \mathbb{F}^{2M_B+8M_O+5M_F+9}$, the encryption of W' under pk^{view}
- $N^{\text{wallet-spend}}(W) \in \mathbb{F}$
- $N^{\text{wallet-match}}(W) \in \mathbb{F}$
- $N^{\text{note-redeem}}(N) \in \mathbb{F}$
- $R_{\text{global}} \in \mathbb{F}$, the current root of the commitment Merkle tree.
- $\omega \in \mathbb{F}$, the flag that designates if the note being settled arose from a match.

I KNOW PRIVATE VARIABLES

- $W = (B, O, F, K, r)$, the old wallet.
- $W' = (B', O', F', K', r')$, the new wallet.
- $\text{Open}(C(W), R_{\text{global}})$, a Merkle proof that $C(W)$ is inserted into the commitment tree.
- $\text{sk}^{\text{settle}} \in \mathbb{F}$
- $N = \left((\hat{m}_1, \hat{v}_1, \hat{d}_1), (\hat{m}_2, \hat{v}_2, \hat{d}_2), (\hat{m}_3, \hat{v}_3, \hat{d}_3), \omega, \hat{r} \right)$, a note.
- $E_N = E_{\text{pk}^{\text{settle}}}(N)$, an encrypted note.
- $O(E_N, R_{\text{global}})$, a Merkle proof that the encryption of N is inserted into the commitment tree.

SUCH THAT

- $C(W)$ is correctly computed.
- $C(W')$ is correctly computed.
- $E_{\text{pk}^{\text{view}}}(W')$ is correctly computed.
- $E_{\text{pk}^{\text{settle}}}(N)$ is correctly computed.
- $N^{\text{wallet-spend}}(W)$ is correctly computed.
- $N^{\text{wallet-match}}(W)$ is correctly computed.
- $N^{\text{note-redeem}}(N)$ is correctly computed.
- $\text{Open}(C(W), R_{\text{global}})$ is a valid Merkle proof.
- $\text{Open}(E_N, R_{\text{global}})$ is a valid Merkle proof.
- $K' = K$
- $r' = r + 2$
- $\text{pk}^{\text{settle}}$ is the valid public key corresponding to $\text{sk}^{\text{settle}}$
- For all balances $b'_i = (m'_i, v'_i) \in B'$:
 - Either $m'_i = 0$, or m'_i is unique in the list of all mints of B' . (i.e., no duplicate mints are allowed).
 - v'_i is equal to

$$\sum_{j \in [M_B] \text{ s.t. } m_j = m'_i} v_j$$

plus

$$\left(\mathbf{1}_{m'_i = \tilde{m}_1 \wedge \tilde{d}_1 = 0} - \mathbf{1}_{m'_i = \tilde{m}_1 \wedge \tilde{d}_1 = 1} \right) \tilde{v}_1$$

plus

$$\left(\mathbf{1}_{m'_i = \tilde{m}_2 \wedge \tilde{d}_2 = 0} - \mathbf{1}_{m'_i = \tilde{m}_2 \wedge \tilde{d}_2 = 1} \right) \tilde{v}_2$$

plus

$$\left(\mathbf{1}_{m'_i = \tilde{m}_3 \wedge \tilde{d}_3 = 0} - \mathbf{1}_{m'_i = \tilde{m}_3 \wedge \tilde{d}_3 = 1} \right) \tilde{v}_3$$

(i.e., balances are unchanged, except for an increase or decrease according to N)

- For all orders $o'_i = (\lambda'_i, m'_{1_i}, m'_{2_i}, s'_i, p'_i, a'_i, \alpha'_i, \tau'_i) \in O'$:
 - $\lambda'_i = \lambda_i, m'_{1_i} = m_{1_i}, m'_{2_i} = m_{2_i}, s'_i = s_i, p'_i = p_i, \alpha'_i = \alpha_i, \tau'_i = \tau_i$
 - a'_i is equal to

$$\sum_{j \in [M_O] \text{ s.t. } m_{1_j} = m'_{1_i} \wedge m_{2_j} = m'_{2_i}} a_j$$

minus

$$\mathbf{1}_{\omega=1 \wedge m'_{1_i} = \hat{m}_1 \wedge m'_{2_i} = \hat{m}_2} \cdot \hat{v}_1$$

(i.e., the orders are unchanged, except for a decrease of the amount corresponding to the matched value)